

Building a Jupyter Kernel

How Python can bring another language to the notebook

Using case study of IForth - a kernel for Forth language

Jupyter Notebook Examples - Python and Forth




The screenshot shows a Jupyter Notebook titled "demo_forth_notebook" with a "Trusted" status. The interface includes a menu bar (File, Edit, View, Run, Kernel, Settings, Help) and a toolbar with icons for file operations and code execution. The notebook content consists of two code cells. The first cell contains the Python code `print("Hello World!")`, which has been executed, resulting in the output "Hello World!". The second cell contains the Python code `2 + 3`, which has been executed, resulting in the output `5`.

```
[1]: print("Hello World!")
Hello World!

[2]: 2 + 3

[2]: 5
```



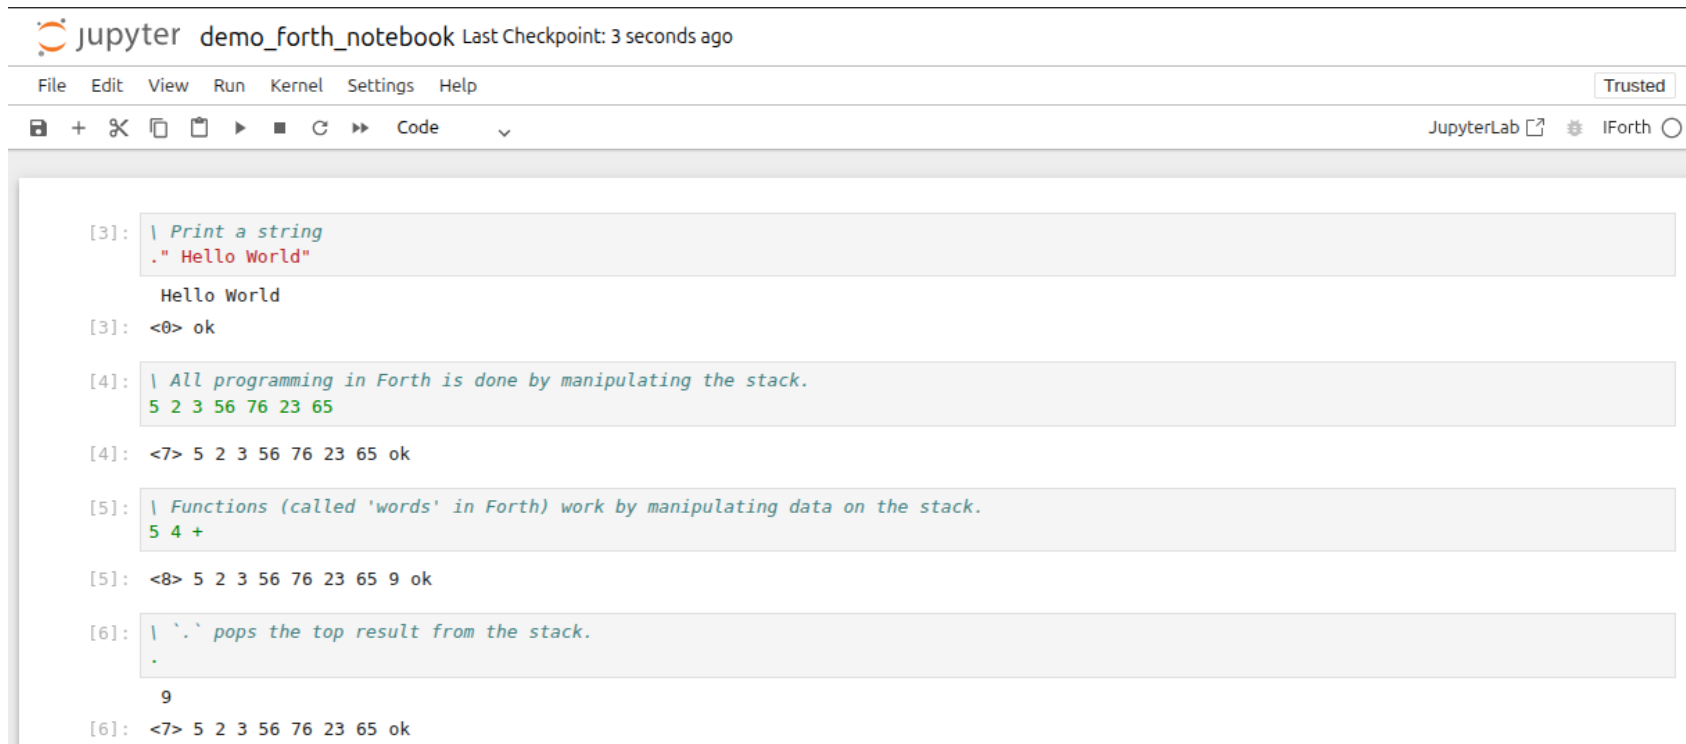
The screenshot shows a Jupyter Notebook titled "demo_forth_notebook" with a "Trusted" status. The interface includes a menu bar (File, Edit, View, Run, Kernel, Settings, Help) and a toolbar with icons for file operations and code execution. The notebook content consists of three code cells. The first cell contains the Forth code `." Hello World!"`, which has been executed, resulting in the output "Hello World!". The second cell contains the Forth code `1 2 3 +`, which has been executed, resulting in the output `<0> ok`. The third cell contains the Forth code `clearstack`, which has been executed, resulting in the output `<0> ok`.

```
[9]: ." Hello World!"
Hello World!
[9]: <0> ok

[6]: 1 2 3 +
[6]: <2> 1 5 ok

[8]: clearstack
[8]: <0> ok
```

Forth in 30 seconds



The screenshot shows a JupyterLab interface with a notebook titled "demo_forth_notebook". The top bar includes the Jupyter logo, the notebook name, and a "Last Checkpoint: 3 seconds ago" status. Below the top bar is a menu bar with "File", "Edit", "View", "Run", "Kernel", "Settings", and "Help". To the right of the menu bar is a "Trusted" button. Below the menu bar is a toolbar with icons for file operations, a "Code" button, and a "JupyterLab" button. The notebook content consists of six code cells, each with a prompt "[3]:", "[4]:", "[5]:", or "[6]:". Each cell contains a Forth code snippet and its output.

```
[3]: | Print a string
    ." Hello World"
    Hello World
[3]: <0> ok


[4]: | All programming in Forth is done by manipulating the stack.
    5 2 3 56 76 23 65
[4]: <7> 5 2 3 56 76 23 65 ok

[5]: | Functions (called 'words' in Forth) work by manipulating data on the stack.
    5 4 +
[5]: <8> 5 2 3 56 76 23 65 9 ok

[6]: | `.` pops the top result from the stack.
    .
    9
[6]: <7> 5 2 3 56 76 23 65 ok
```

See more syntax examples: <https://learnxinyminutes.com/forth>

The message flow



The Jupyter Server communicates with the kernel over 5 ZeroMQ channels:

- Shell: Main request/reply loop, frontend sends code to run here.
- IOPub: Broadcast channel. Kernel can publish stdout, stderr, HTML, etc.
- Stdin: Kernel can request user input.
- Control: High-priority commands (eg. Shutdown, Interrupt) that bypass execution queue.
- Heartbeat: Simple ping/pong socket, check kernel is still alive.

Source: <https://www.datahaskell.org/blog/2025/11/25/a-tale-of-two-kernels.html#the-jupyter-kernel-architecture>

Python is the glue

Jupyter speaks a protocol; Python connects it to Forth.

- Jupyter's Kernel base class provides the integration point
- Python starts and owns a long-lived GForth subprocess
- stdin / stdout / stderr connect Python to the interpreter
- asyncio lets the kernel stream output without blocking Jupyter

From a Python cell to GForth

A cell execution becomes a small Python pipeline:

Jupyter → **do_execute(code)** → **GForth stdin** → **GForth stdout/stderr** → Jupyter

- `do_execute()` receives the cell's source code
- GForth is kept alive between cells, so state persists
- Python reads output until the GForth prompt
- `stdout/stderr` are sent back as Jupyter stream messages

The Python side: two kinds of output

Printed output	Expression / stack value
answer_text(...)	answer_expression_value(...)
stream message	execute_result message

This separation is what makes a non-Python interpreter feel native inside a notebook.

In IForth, Python implements the Jupyter-facing layer; GForth remains the Forth runtime.

Subclass Jupyter Kernel

```
class IForth(Kernel):
    """Jupyter kernel for Forth language."""
    implementation = "forth"
    implementation_version = __version__

    gforth = Gforth()  # spawns subprocess: long-lived gforth interpreter

    def answer_text(self, text: str, stream: Literal["stdout", "stderr"]) -> None:
        """Send text response to Jupyter cell."""
        logger.info("Answering to Jupyter: %s", text)
        self.send_response(
            self.iopub_socket, "stream", {"name": stream, "text": text}
        )

    def answer_expression_value(self, expression_value: str) -> None:
        self.send_response(self.iopub_socket, "execute_result", {
            "execution_count": self.execution_count,
            "data": {
                "text/plain": expression_value,
            },
            "metadata": {}
        })

    @override
    async def do_execute(  # do_execute() can also be synchronous
        self,
        code: str,
        silent: bool,
        store_history: bool = True,
        user_expressions = None,
        allow_stdin = None,
        *,
        cell_meta = None,
        cell_id = None,
    ) -> dict[str, Any]:
        """Execute Forth code."""
        logger.info("do_execute (silent=%s): %s", silent, code)
        stack_output = await self.gforth.exec(code, self.answer_text)
        if stack_output is not None:
            self.answer_expression_value(stack_output)
        return {
            'status': 'ok',
            'execution_count': self.execution_count,  # Jupyter's tracked cell count
            'payload': [],
            'user_expressions': {}
        }
```


Register the custom kernel with Jupyter

```
# dynamically generate kernel.json so that correct Python path is used
kernel_spec = {
    "argv": [
        # path to python (of the venv where forth kernel is installed)
        # allows it to run even if Jupyter is installed in a different venv
        sys.executable,
        "-m",
        "forth_kernel",
        "-f",
        "{connection_file}"
    ],
    "display_name": "IForth",
    "language": "Forth",
    "codemirror_mode": "text",
    "interrupt_mode": "message",
    "name": "Forth"
}

script_dir = Path(__file__).parent.resolve()
logger.info('Script Dir: %s', script_dir)
with (script_dir / 'kernel.json').open('w') as f:
    json.dump(kernel_spec, f, indent=4)

try:
    install_kernel_spec(str(script_dir), 'forth', replace=True, user=args.user)
    logger.info('Successfully installed jupyter kernel for Forth.')
except PermissionError as e:
    logger.error(
        'Failed to install jupyter kernel for Forth: %s\n' +
        'Try installing as user instead of root: python -m forth_kernel.self_install --user',
        e
    )
except Exception as e:
    logger.error('Failed to install jupyter kernel for Forth: %s', e)
```

Automated Tests using module jupyter_kernel_test

```
class MyKernelTests(jupyter_kernel_test.KernelTests):
    kernel_name = "forth"
    language_name = "forth"
    file_extension = ".4th"

    code_hello_world = '." hello, world"'
    code_stderr = 'clearstack .' # stack underflow error
    code_execute_result = [
        {'code': '1 2 3 .', 'result': '<2> 1 2 ok'},
    ]

    def setUp(self):
        self.flush_channels()

    def test_multiline_cell_is_prompt(self):
        """Regression: each line used to cost a full read timeout."""
        start = time.monotonic()
        reply, _ = self.execute_helper(code='1 .\n2 .\n3 .\n4 .\n5 .')
        self.assertEqual(reply['content']['status'], 'ok')
        self.assertLess(time.monotonic() - start, 5)

    def test_long_output_not_truncated(self):
        """Regression: output was cut off when it exceeded the read timeout."""
        _, msgs = self.execute_helper(code=': countup 200 0 do i . loop ; countup')
        out = ''.join(m['content']['text'] for m in msgs)
        if m['msg_type'] == 'stream'
            and m['content']['name'] == 'stdout')
        self.assertIn('199', out)

if __name__ == "__main__":
    import unittest
    unittest.main()
```

Check out IForth for an example Jupyter kernel implementation.

And maybe try Forth - it's an interesting language :)

Repo: <https://github.com/sohang3112/iforth>

Thank You!

PS: This presentation was made using <https://docs.racket-lang.org/slideshow/>

Questions?